

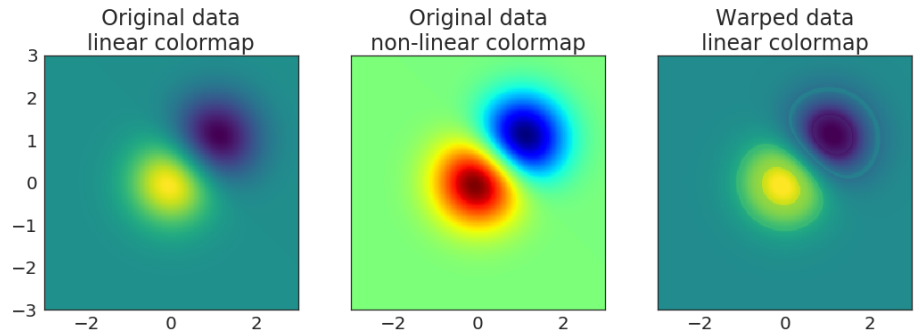
Stay away from the jet (rainbow color map)

Luminance spikes in center



<http://jakevdp.github.io/blog/2014/10/16/how-bad-is-your-colormap/>


Warps your perception of the visualized data, seeing non-linearities when there aren't any



<https://predictablynoisy.com/makeitpop-intro>

Instead, perhaps use these!

Sequential:
viridis or plasma



Diverging:
RdBu_r (reversed) or PrGn_r



Categorical:
Accent/Dark2



Reading CSV

Read straight from URL! So cool!


```
import pandas as pd

url = 'https://mesonet.agron.iastate.edu/cgi-bin/request/asos'
df = pd.read_csv(url, index_col='valid', parse_dates=True)
print(df.head())  # print first 5
```

		station	tmpf
valid			
2018-01-01	00:00:00	ALO	M
2018-01-01	00:05:00	ALO	M
2018-01-01	00:10:00	ALO	M
2018-01-01	00:15:00	ALO	M
2018-01-01	00:20:00	ALO	M


Plot diurnal cycle

```
df.groupby(df.index.hour).mean().plot()
```



Get an interactive plot

```
import holoplot.pandas  
df.groupby(df.index.hour).mean().holoplot()
```



Pandas can also read unformatted tables!

```
RMM values up to "real time". For the last few days, ACCESS analyses are used instead of NCEP
year, month, day, RMM1, RMM2, phase, amplitude. Missing Value= 1.E36 or 999
1974      6      1  1.6344700      1.2030400      5  2.0294800      Final_value: __OLR & NCEP_winds
1974      6      2  1.6028900      1.0151200      5  1.8972900      Final_value: __OLR & NCEP_winds
1974      6      3  1.5162500      1.0855100      5  1.8647600      Final_value: __OLR & NCEP_winds
1974      6      4  1.5098100      1.0357300      5  1.8309200      Final_value: __OLR & NCEP_winds
1974      6      5  1.5590600      1.3051800      5  2.0332601      Final_value: __OLR & NCEP_winds
1974      6      6  1.2062600      1.6288900      6  2.0269001      Final_value: __OLR & NCEP_winds
1974      6      7  0.61110097      1.7224801      6  1.8276700      Final_value: __OLR & NCEP_winds
1974      6      8  0.32639500      1.7781800      6  1.8078901      Final_value: __OLR & NCEP_winds
```

```
df = pd.read_table('rmm.74toRealtime.txt')
df.head()
```

RMM values up to "real time". For the last few days, ACCESS analyses are used instead of NCEP

	year, month, day, RMM1, RMM2, phase, amplitude...	Final_value: __OLR & NCEP_winds
0		
1		
2		
3		
4		

year	month	day	RMM1	RMM2	phase	amplitude	Missing Value= 1.E36 or 999			
1974	6	1	1.6344700	1.2030400	5	2.0294800		Final_value: __OLR_ & NCEP_winds		
1974	6	2	1.6028900	1.0151200	5	1.8972900		Final_value: __OLR_ & NCEP_winds		
1974	6	3	1.5162500	1.0855100	5	1.8647600		Final_value: __OLR_ & NCEP_winds		
1974	6	4	1.5098100	1.0357300	5	1.8309200		Final_value: __OLR_ & NCEP_winds		
1974	6	5	1.5590600	1.3051800	5	2.0332601		Final_value: __OLR_ & NCEP_winds		
1974	6	6	1.2062600	1.6288900	6	2.0269001		Final_value: __OLR_ & NCEP_winds		
1974	6	7	0.61110097	1.7224801	6	1.8276700		Final_value: __OLR_ & NCEP_winds		
1974	6	8	0.32639500	1.7781800	6	1.8078901		Final_value: __OLR_ & NCEP_winds		

```
df = pd.read_csv('rmm.74toRealtime.txt',
                 skiprows=1,
                 delimiter='\s+',
                 parse_dates=[['year', 'month', 'day']],
                 index_col='year_month_day')
df.columns = df.columns.str.replace(',', '').str.replace('.', '')
df = df[df.columns[:4]]
df.index.name = 'datetime'
df.head()
```

	RMM1	RMM2	phase	amplitude
datetime				
1974-06-01	1.63447	1.20304	5	2.02948
1974-06-02	1.60289	1.01512	5	1.89729
1974-06-03	1.51625	1.08551	5	1.86476
1974-06-04	1.50981	1.03573	5	1.83092
1974-06-05	1.55906	1.30518	5	2.03326

year, month, day,	RMM1,	RMM2,	phase,	amplitude.	Missing Value=	1.E36 or 999		
1974	6	1	1.6344700	1.2030400	5	2.0294800	Final_value: __OLR_ & NCEP_winds	
1974	6	2	1.6028900	1.0151200	5	1.8972900	Final_value: __OLR_ & NCEP_winds	
1974	6	3	1.5162500	1.0855100	5	1.8647600	Final_value: __OLR_ & NCEP_winds	
1974	6	4	1.5098100	1.0357300	5	1.8309200	Final_value: __OLR_ & NCEP_winds	
1974	6	5	1.5590600	1.3051800	5	2.0332601	Final_value: __OLR_ & NCEP_winds	
1974	6	6	1.2062600	1.6288900	6	2.0269001	Final_value: __OLR_ & NCEP_winds	
1974	6	7	0.61110097	1.7224801	6	1.8276700	Final_value: __OLR_ & NCEP_winds	
1974	6	8	0.32639500	1.7781800	6	1.8078901	Final_value: __OLR_ & NCEP_winds	

```
df = pd.read_csv('rmm.74toRealtime.txt',
                 skiprows=1,
                 delimiter='\s+',
                 parse_dates=[['year,', 'month,', 'day,']],
                 index_col='year_month_day,')
df.columns = df.columns.str.replace(',', '').str.replace('.', '')
df = df[df.columns[:4]]
df.index.name = 'datetime'
df.head()
```

	RMM1	RMM2	phase	amplitude
datetime				
1974-06-01	1.63447	1.20304	5	2.02948
1974-06-02	1.60289	1.01512	5	1.89729
1974-06-03	1.51625	1.08551	5	1.86476
1974-06-04	1.50981	1.03573	5	1.83092
1974-06-05	1.55906	1.30518	5	2.03326

year	month	day	RMM1	RMM2	phase	amplitude	Missing Value= 1.E36 or 999		
1974	6	1	1.6344700	1.2030400	5	2.0294800		Final_value: __OLR_ & NCEP_winds	
1974	6	2	1.6028900	1.0151200	5	1.8972900		Final_value: __OLR_ & NCEP_winds	
1974	6	3	1.5162500	1.0855100	5	1.8647600		Final_value: __OLR_ & NCEP_winds	
1974	6	4	1.5098100	1.0357300	5	1.8309200		Final_value: __OLR_ & NCEP_winds	
1974	6	5	1.5590600	1.3051800	5	2.0332601		Final_value: __OLR_ & NCEP_winds	
1974	6	6	1.2062600	1.6288900	6	2.0269001		Final_value: __OLR_ & NCEP_winds	
1974	6	7	0.61110097	1.7224801	6	1.8276700		Final_value: __OLR_ & NCEP_winds	
1974	6	8	0.32639500	1.7781800	6	1.8078901		Final_value: __OLR_ & NCEP_winds	

```
df = pd.read_csv('rmm.74toRealtime.txt',
                 skiprows=1,
                 delimiter='\s+',
                 parse_dates=[['year', 'month', 'day']],
                 index_col='year_month_day')
df.columns = df.columns.str.replace(',', '').str.replace('.', '')
df = df[df.columns[:4]]
df.index.name = 'datetime'
df.head()
```

	RMM1	RMM2	phase	amplitude
datetime				
1974-06-01	1.63447	1.20304	5	2.02948
1974-06-02	1.60289	1.01512	5	1.89729
1974-06-03	1.51625	1.08551	5	1.86476
1974-06-04	1.50981	1.03573	5	1.83092
1974-06-05	1.55906	1.30518	5	2.03326

RMM values up to "real time". For the last few days, ACCESS analyses are used instead of NCEP
year, month, day, RMM1, RMM2, phase, amplitude. Missing Value= 1.E36 or 999

1974	6	1	1.6344700	1.2030400	5	2.0294800	Final_value: __OLR_ & NCEP_winds
1974	6	2	1.6028900	1.0151200	5	1.8972900	Final_value: __OLR_ & NCEP_winds
1974	6	3	1.5162500	1.0855100	5	1.8647600	Final_value: __OLR_ & NCEP_winds
1974	6	4	1.5098100	1.0357300	5	1.8309200	Final_value: __OLR_ & NCEP_winds
1974	6	5	1.5590600	1.3051800	5	2.0332601	Final_value: __OLR_ & NCEP_winds
1974	6	6	1.2062600	1.6288900	6	2.0269001	Final_value: __OLR_ & NCEP_winds
1974	6	7	0.61110097	1.7224801	6	1.8276700	Final_value: __OLR_ & NCEP_winds
1974	6	8	0.32639500	1.7781800	6	1.8078901	Final_value: __OLR_ & NCEP_winds

```
df = pd.read_csv('rmm.74toRealtime.txt',
                 skiprows=1,
                 delimiter='\s+',
                 parse_dates=[['year','month','day']],
                 index_col='year_month_day')
df.columns = df.columns.str.replace(',','').str.replace('.', '')
df = df[df.columns[:4]]
df.index.name = 'datetime'
df.head()
```

Beautiful, useable output in a just few lines with no “for loops”!

```
df = pd.read_csv('rmm.74toRealtime.txt',
                 skiprows=1,
                 delimiter='\s+',
                 parse_dates=[['year','month','day']],
                 index_col='year_month_day')
df.columns = df.columns.str.replace(',','').str.replace('.', '')
df = df[df.columns[:4]]
df.index.name = 'datetime'
df.head()
```

	RMM1	RMM2	phase	amplitude
datetime				
1974-06-01	1.63447	1.20304	5	2.02948
1974-06-02	1.60289	1.01512	5	1.89729
1974-06-03	1.51625	1.08551	5	1.86476
1974-06-04	1.50981	1.03573	5	1.83092
1974-06-05	1.55906	1.30518	5	2.03326

Reading data with higher dimensions: netCDF

```
import xarray as xr
ds = xr.open_mfdataset('air.sig995.20*.nc')
print(ds)
```

```
<xarray.Dataset>
Dimensions:      (lat: 73, lon: 144, nbnds: 2, time: 886)
Coordinates:
  * lat          (lat) float32 90.0 87.5 85.0 82.5 80.0 77.5 75.0 72.5 70.0 ...
  * lon          (lon) float32 0.0 2.5 5.0 7.5 10.0 12.5 15.0 17.5 20.0 22.5 ...
  * time         (time) datetime64[ns] 2016-01-01 2016-01-02 2016-01-03 ...
Dimensions without coordinates: nbnds
Data variables:
    air          (time, lat, lon) float32 dask.array<shape=(886, 73, 144), chunksize=(366, 73, 144)>
    time_bnds    (time, nbnds) float64 dask.array<shape=(886, 2), chunksize=(366, 2)>
Attributes:
    Conventions: COARDS
    title:       mean daily NMC reanalysis (2014)
    history:     created 2013/12 by Hoop (netCDF2.3)
    description: Data is from NMC initialized reanalysis\n(4x/day).  These...
    platform:    Model
    References:   http://www.esrl.noaa.gov/psd/data/gridded/data.ncep.reana...
    dataset_title: NCEP-NCAR Reanalysis 1
```

Get the monthly “climatology” (only two years here)

```
ds_air = ds['air']
ds_air_avg = ds_air.groupby('time.month').mean('time') # two year monthly mean (ten plus years = monthly climatology)
ds_air_avg
```

```
<xarray.DataArray 'air' (month: 12, lat: 73, lon: 144)>
dask.array<shape=(12, 73, 144), dtype=float32, chunksize=(1, 73, 144)>
Coordinates:
  * lat      (lat) float32 90.0 87.5 85.0 82.5 80.0 77.5 75.0 72.5 70.0 67.5 ...
  * lon      (lon) float32 0.0 2.5 5.0 7.5 10.0 12.5 15.0 17.5 20.0 22.5 ...
  * month    (month) int64 1 2 3 4 5 6 7 8 9 10 11 12
```

Similar to NCL's clmMonTLL

Get the daily anomalies

```
ds_air_anom = ds_air.groupby('time.month') - ds_air_avg
```

Similar to NCL's calcMonAnomTLL

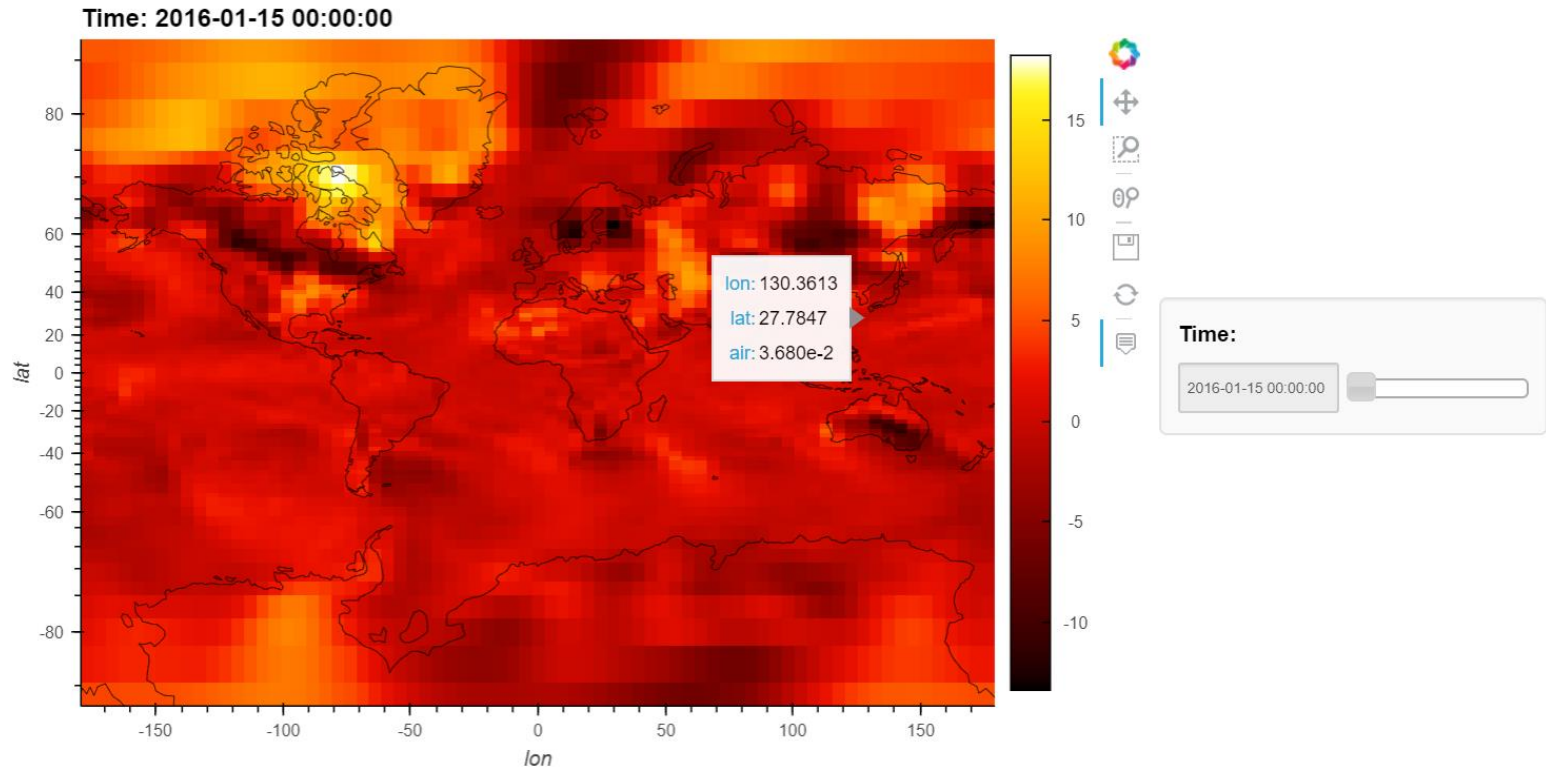
Get interactive map easily!

```
In [17]: import geoviews as gv
import holovplot.xarray

air = ds_air_anom.holovplot('lon', 'lat',
                           groupby='time',
                           geo=True,
                           width=725,
                           height=500) * gv.feature.coastline()

air
```

Out[17]:



Reading GRIB

```
import xarray as xr

ds = xr.open_dataset('prate.2018050100.01.CFSv2.anom.avrg.1x1.grb', engine='pynio')
print(ds)
```

<xarray.Dataset>

Dimensions: (forecast_time0: 8, lat_3: 181, lon_3: 360)

Coordinates:

- * lat_3 (lat_3) float32 90.0 89.0 88.0 87.0 86.0 85.0 84.0 ...
- * lon_3 (lon_3) float32 0.0 1.0 2.0 3.0 4.0 5.0 6.0 7.0 8.0 ...
- * forecast_time0 (forecast_time0) int32 2 3 4 5 6 7 8 9

Data variables:

PRATE_3_SFC_ave1m (forecast_time0, lat_3, lon_3) float32 ...

But those data has metadata! What about models that output plain binary?

```
import numpy as np

ntime = 1
nlat = 181
nlon = 360
fn = 'grdtmp_sfc_0000.0_0000.0_glob360x181_2017082012_01020000_fcstfld'
data = np.fromfile(fn, dtype='>f').reshape((ntime, nlat, nlon))
print(data)
```

```
[[[217.52127 217.52124 217.52122 ... 217.52129 217.52129 217.52129]
  [216.6976 216.66498 216.67284 ... 216.69624 216.68231 216.67058]
  [215.18414 215.0505 215.03452 ... 215.38925 215.3056 215.22244]
  ...
  [271.52542 271.51984 271.493 ... 271.4726 271.44864 271.48325]
  [271.51227 271.50754 271.49915 ... 271.5336 271.5274 271.51855]
  [272.10092 272.10092 272.10092 ... 272.10092 272.10092 272.10092]]]
```


Processing plain binary

Make xr.DataArray

```
import pandas as pd

LAT_RANGE = (-90, 90)
LON_RANGE = (0, 360)

dlon = LON_RANGE[1] / nlon
lon = np.arange(LON_RANGE[0], LON_RANGE[1], dlon)

dlat = (LAT_RANGE[1] - LAT_RANGE[0]) / (nlat - 1)
lat = np.arange(LAT_RANGE[0], LAT_RANGE[1] + dlat, dlat)

time = pd.to_datetime(fn.split('_')[5], format='%Y%m%d%H')

ds = xr.DataArray(data,
                  coords={'time': [time], 'lat': lat, 'lon': lon},
                  dims=('time', 'lat', 'lon'))

print(ds)
```

<xarray.DataArray (time: 1, lat: 181, lon: 360)>
array([[[217.52127, 217.52124, ..., 217.52129, 217.52129],
 [216.6976 , 216.66498, ..., 216.68231, 216.67058],
 ...,
 [271.51227, 271.50754, ..., 271.5274 , 271.51855],
 [272.10092, 272.10092, ..., 272.10092, 272.10092]]], dtype=float32)

Coordinates:
* time (time) datetime64[ns] 2017-08-20T12:00:00
* lat (lat) float64 -90.0 -89.0 -88.0 -87.0 -86.0 -85.0 -84.0 -83.0 ...
* lon (lon) float64 0.0 1.0 2.0 3.0 4.0 5.0 6.0 7.0 8.0 9.0 10.0 11.0 ...

Reading Shapefiles

```
import geopandas as gpd

gdf = gpd.read_file('cb_2017_us_county_500k.shp')
gdf.head()
```

	STATEFP	COUNTYFP	COUNTYNS	AFFGEOID	GEOID	NAME	LSAD	ALAND	AWATER	
0	01	005	00161528	0500000US01005	01005	Barbour	06	2292144656	50538698	POLYGON ((-
1	01	023	00161537	0500000US01023	01023	Choctaw	06	2365869837	19144469	POLYGON ((-
2	01	035	00161543	0500000US01035	01035	Conecuh	06	2201948618	6643480	POLYGON ((-
3	01	051	00161551	0500000US01051	01051	Elmore	06	1601762124	99965171	POLYGON ((-
4	01	065	00161558	0500000US01065	01065	Hale	06	1667907107	32423356	POLYGON (

Processing Shapefiles

Easily figure out the top 5 states with the highest population/sq mi

```
gdf['pop_per_mi_sq'] = (gdf['DP0010001'] / gdf['ALAND10'] / 3.86102e-7) # population per square mile
gdf[['NAME10', 'pop_per_mi_sq']].sort_values('pop_per_mi_sq', ascending=False).head() # sort and keep top 5
```

	NAME10	pop_per_mi_sq
--	--------	---------------

46	District of Columbia	9856.491855
----	----------------------	-------------

23	New Jersey	1195.490197
----	------------	-------------

7	Puerto Rico	1088.211268
---	-------------	-------------

5	Rhode Island	1018.139860
---	--------------	-------------

30	Massachusetts	839.433705
----	---------------	------------

HoloViews

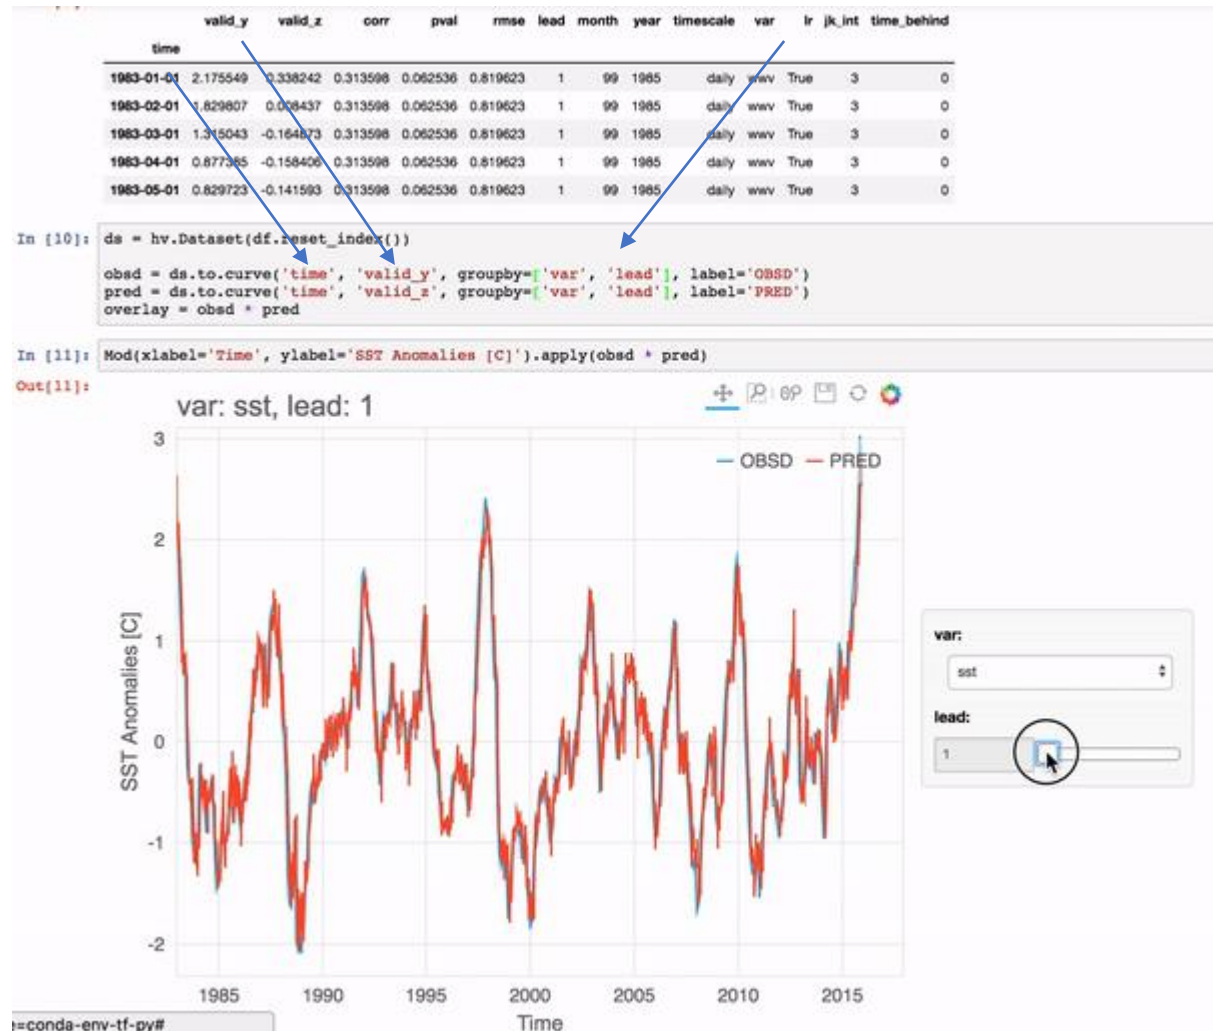
Have a DataFrame

Wrap hv.Dataset

Convert to curve by
providing column
names to x, y, groupby

Multiply two curves to
stack them

Have awesome
clickable animations!



Cartopy

```
import matplotlib.pyplot as plt
import cartopy.crs as ccrs
import cartopy.feature as cfeature

fig = plt.figure(figsize=(12, 8))
projection = ccrs.PlateCarree()
ax = plt.axes(projection=projection)
im = ax.contourf(da.lon, da.lat, da[0].values, transform=projection, cmap='RdBu_r')
_ = plt.colorbar(im, orientation='horizontal', shrink=0.5, pad=0.05)
_ = ax.gridlines(crs=ccrs.PlateCarree(), draw_labels=True)
_ = ax.coastlines()
plt.grid()
```

